

# Homework 9

Yuxin Sun

---


## Results

With

```
batch_size=16,  
epoch_size=15,  
lr=1e-4,  
betas=(0.9, 0.99),  
measure_rate=20
```

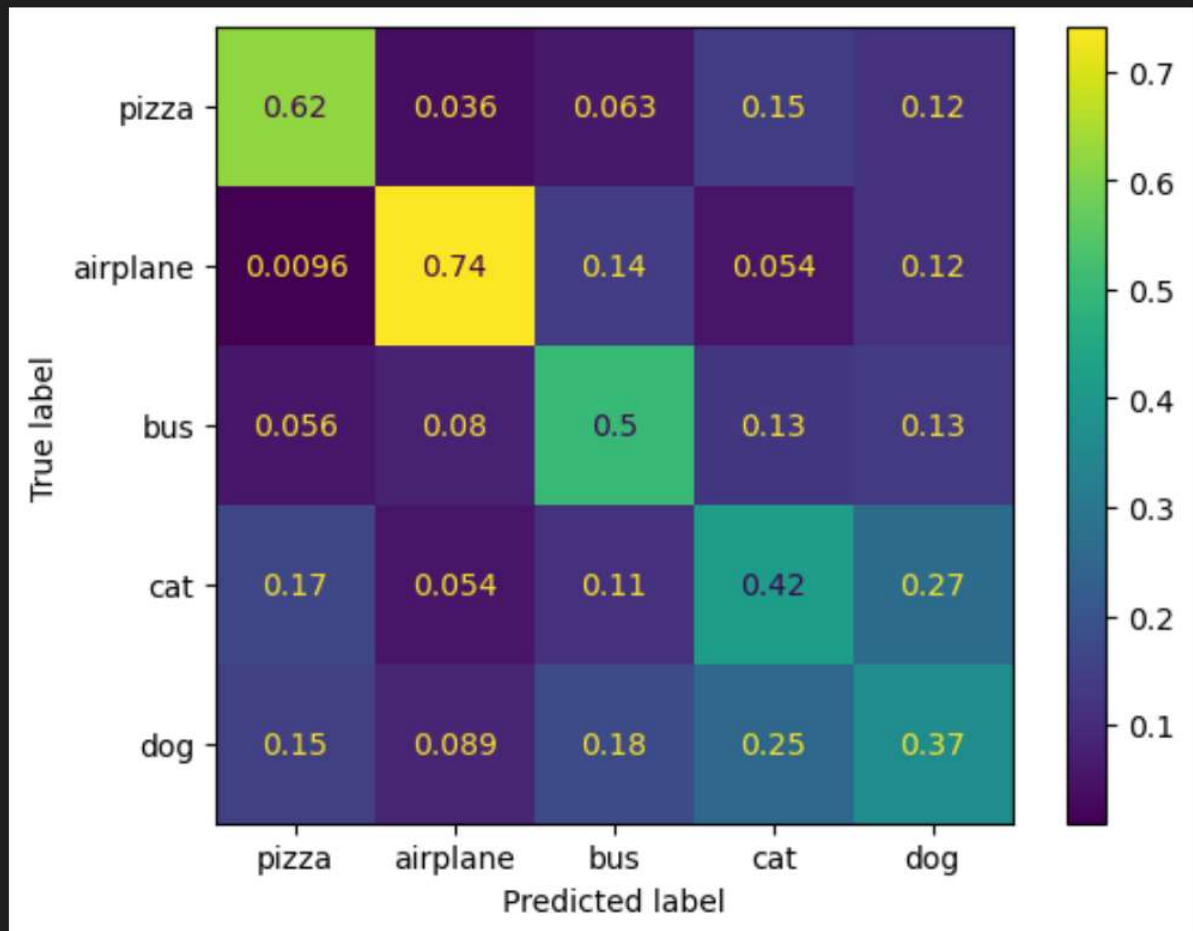
I get

Accuracy: 53.20%



(Figure 1: Loss)

```
tensor([[324., 16., 40., 75., 45.],
        [ 5., 332., 90., 28., 45.],
        [29., 36., 318., 66., 51.],
        [86., 24., 72., 215., 103.],
        [77., 40., 111., 131., 141.]])
```



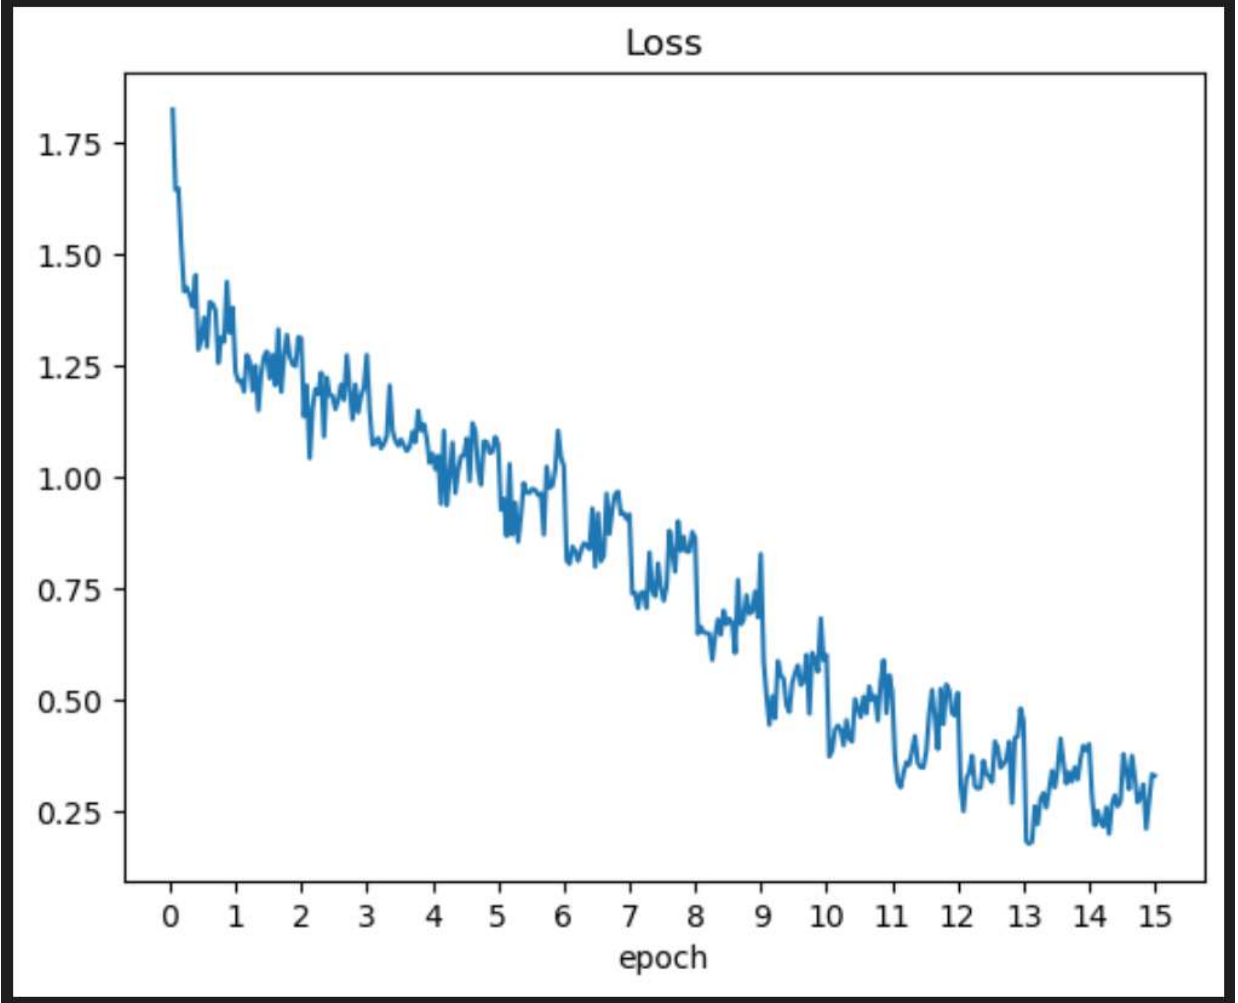
(Figure 2: Confusion Matrix)

## Using `torch.einsum()`

The accuracy should be the same as before.


The smaller accuracy is probably due to no `bias` used in this implementation.

Accuracy: 46.76%



(Figure 3: Loss)


```
tensor([[302.,  8., 27., 83., 80.],  
       [ 23., 283., 77., 32., 85.],  
       [ 35., 60., 228., 61., 116.],  
       [ 98., 16., 65., 183., 138.],  
       [ 89., 49., 68., 121., 173.]])
```



(Figure 4: Confusion Matrix)

## Datasets

Same as hw4...



## Codes

### Multi-head Attention Using `torch.einsum()`:

```
class BasicEncoder(nn.Module):
    def __init__():
        ...
        # replace SelfAttention with SelfAttention2
        self.self_attention_layer = SelfAttention2(
            max_seq_length, embedding_size, num_atten_heads) # (A)
        ...

class SelfAttention2(nn.Module):
    def __init__(self, max_seq_length, embedding_size, num_atten_heads):
        super().__init__()
        self.max_seq_length = max_seq_length
        self.embedding_size = embedding_size
        self.num_atten_heads = num_atten_heads
```

```

        self.qkv_size = self.embedding_size // num_attn_heads
        # we just need one tensor:
        self.wqkv_tensor = nn.Parameter(
            torch.randn(num_attn_heads, 3, max_seq_length, self.qkv_size,
max_seq_length, self.qkv_size),
            requires_grad=True
        ).to(device)
        self.coeff =
1.0/torch.sqrt(torch.tensor(self.qkv_size).float()).to(device)

    def forward(self, sentence_tensor):
        b, s, d = sentence_tensor.shape
        x = sentence_tensor.view(b, s, self.num_attn_heads, self.qkv_size)
        # b: batch index,
        # s: sequence index, si sj --> i j
        # h: attention-head index,
        # q: S_{qkv}
        q, k, v = tuple(torch.einsum('bshq,hksqSQ->kbhSQ', x,
self.wqkv_tensor))
        QK_dot_prod = torch.einsum('bhiq,bhjq->bhij', q, k)
        rowwise_softmax_normalizations = torch.softmax(QK_dot_prod, dim=-1)
        z = torch.einsum('bhsj,bhjq->bshq', rowwise_softmax_normalizations, v)
        z = self.coeff * z.reshape(b,s,d)
        return z

```

## My ViT Net

```

class ViT(nn.Module):
    def __init__(self, *, max_seq_length=17, embedding_size=64, encoders=2,
num_attn_heads=4,
                device="cpu"):
        super().__init__()
        self.encoder = MasterEncoder(max_seq_length, embedding_size, encoders,
num_attn_heads)
        # transform sentences to embeddings:
        self.embedding = nn.Sequential(
            nn.Conv2d(3, embedding_size, 16, stride=16, padding=0),
            nn.Flatten(start_dim=2, end_dim=-1)
        )
        self.class_head = nn.Parameter(torch.randn(embedding_size,
device=device), requires_grad=True)
        self.position = nn.Parameter(2*torch.rand(max_seq_length,
embedding_size, device=device)-1,requires_grad=True)
        # MLP with only one hidden layer is enough
        self.mlp = nn.Sequential(
            nn.Flatten(),
            nn.Linear(max_seq_length*embedding_size, 128),
            nn.ReLU(),
            nn.Linear(128, 5)
        )

    def forward(self, x):
        # x_out is of (B, 16, D)
        x = self.embedding(x).transpose(1,2)
        # class_head: (B, 1, D)

```

```

class_head = self.class_head.repeat(x.shape[0],1,1)

x = torch.cat((x, class_head), dim=1).add(self.position)
x = self.encoder(x)
out = self.mlp(x)
return out

def train(self, p):
    # p contains all the runtime parameters
    self.device = p.device
    net = self.to(p.device)
    start_time = time.perf_counter()
    criterion = nn.CrossEntropyLoss()
    optimizer = torch.optim.Adam(net.parameters(), lr=p.lr, betas=p.betas)
    loss_record = []
    for epoch in range(1, p.epoch_size+1):
        running_loss = 0.
        epoch_start_time = time.perf_counter()
        for i, data in enumerate(p.train_dataloader):
            img_tensor, label_tensor = data
            img_tensor = img_tensor.to(p.device)
            label_tensor = label_tensor.to(p.device)
            optimizer.zero_grad()
            output = net(img_tensor)
            loss = criterion(output, label_tensor)
            running_loss += loss.item()
            loss.backward()
            optimizer.step()

            j = i+1
            if j % p.measure_rate == 0:
                avg_loss = running_loss/p.measure_rate
                running_loss = 0.
                loss_record.append(avg_loss)
                current_time = time.perf_counter()
                time_elapsed = current_time - epoch_start_time
                print("[epoch:%d/%d] [iter:%4d] elapsed time:%4d secs
loss: %.5f\"
                                %(epoch, p.epoch_size, j, time_elapsed, avg_loss) )

time_stamp = '_' .join(time.ctime().split(" ")[1:4])
saved_model = p.save_dir + "saved_model_" + time_stamp
torch.save(net.state_dict(), saved_model)
print("model.state_dict() saved to ", saved_model)
saved_file = p.save_dir + "saved_file_" + time_stamp
with open(saved_file, 'w') as f:
    for loss in loss_record:
        f.write("%.5f\n" % loss)
    f.flush()
total_time = time.perf_counter() - start_time
print("Total training time: %5d secs"%total_time)
return loss_record, time_stamp

def test(self, p, time_stamp):

```

```

        self.load_state_dict(torch.load(p.save_dir + "saved_model_" +
time_stamp))
        self.device = p.device
        net = self.to(p.device)
        accuracy = 0
        cf_matrix = torch.zeros(5,5)
        with torch.no_grad():
            for i, data in enumerate(p.test_dataloader):
                img_tensor, label_tensor = data
                img_tensor = img_tensor.to(p.device)
                label_tensor = label_tensor.to(p.device)
                output = net(img_tensor)
                predicted_idx = torch.argmax(output, dim=1)
                for label, prediction in zip(label_tensor, predicted_idx):
                    cf_matrix[label][prediction] += 1
                    if label == prediction:
                        accuracy += 1
        accuracy = float(accuracy) / sum(cf_matrix.view(-1))
        print("Accuracy: %.2f%%"%(accuracy*100))
        return cf_matrix

```

## Discussion

Is the network performance better than CNN version?

- My ViT can recognize pizza better than CNN.
- Its performance for the other 4 classes is always worse than CNN.